

Parallel(?) Tempering

Rômulo Cenci

December 13, 2019




So.. Who am I?

- ▶ Universidade Federal de Santa Catarina - Florianópolis, Brazil.
- ▶ Advisor: Lucas Nicolao
- ▶ Research area: Condensed Matter/Statistical Mechanics




My masters project

- ▶ I'm studying a model called GEM- α that could model some colloidal suspensions and ultracold atoms;
- ▶ And that forms some special patterns know as cluster cristals.



What is my problem?


- ▶ We generally want to get out of metastable states;
- ▶ Equilibrate simulation in various temperatures;
- ▶ We don't want to waste time doing bad annealing that may not be in a equilibrium state.



- ▶ So, to solve all that problems we can use the **parallel tempering** technique;

Parallel Tempering

- ▶ In my simulations, the temperature is a fixed variable and energy could fluctuate → Canonical ensemble;
- ▶ So if we look at the energy histograms, we will see an overlap between different temperatures;




- ▶ This overlap could be interpreted as the probability of the configuration at this temperature could be at the other temperature;
- ▶ We can calculate the total probability of two configurations have been changed using the Boltzmann weight;

$$p(i, j) = \min \left(1, \frac{e^{-\beta_i E_j - \beta_j E_i}}{e^{-\beta_i E_i - \beta_j E_j}} \right) = \min \left(1, e^{(E_i - E_j)(\beta_i - \beta_j)} \right)$$

How can we parallelize?

- Just a task farming, so...



```

for(trials=0; trials<300; trials++){
    MPI_Barrier(MPI_COMM_WORLD);
    MPI_Scatter(allT, 1, MPI_DOUBLE, &T, 1, MPI_DOUBLE, 0, comm);

    // update all the ensembles in the same time, one per thread
    for(k=0; k<Kmax; k++){
        update(x, y, T);
    }

    // wait all the evolutions finish
    MPI_Barrier(MPI_COMM_WORLD);

    // measure the energy of each ensemble and send to master
    E = ener(x, y);

    MPI_Gather(&E, 1, MPI_DOUBLE, allE, 1, MPI_DOUBLE, 0, comm);

    // decides who changes temperature with who
    if(pt==0){
        for(i=0; i<PT-1; i++){
            delta=-(1./allT[confT[i+1]]-1./allT[confT[i]])*(allE[confT[i+1]]-allE[confT[i]]);
            if(delta<0 || FRANDOM<exp(-delta)){
                aux=confT[i];
                confT[i]=confT[i+1];
                confT[i+1]=aux;

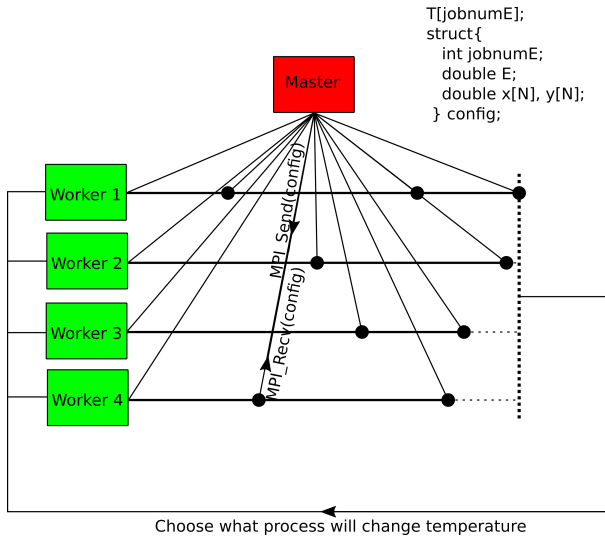
                aux=allT[i+1];
                allT[i+1]=allT[i];
                allT[i]=aux;
            }
        }
    }
}

```

Load Balance

- ▶ Okay, so now we have a truly **parallel** tempering, running at different cores;
- ▶ But to running N_{temps} temperatures we need N_{temps} cores;
- ▶ So, when we have a small cluster, how the ours: 14 quad-core nodes = 48 processing elements, we can just run 48 temperatures;
- ▶ How to simulate more temperatures in this environment without losing much time with one process and almost anything with the other?

Load Balance



```
T[jobnumE];  
struct{  
    int jobnumE;  
    double E;  
    double x[N], y[N];  
} config;
```

```

// master process
if(!rank){
    // initiates sending one process to each rank
    for(jobnum=0; jobnum<Nranks-1; jobnum++){
        configs[jobnum].jobnumE = jobnum;
        MPI_Send(&configs[jobnum], 1, MPI_Config,\
                jobnum+1, flag, comm);
    }

    // still has jobs to do
    for(jobnum=Nranks-1; jobnum<Ntemps; jobnum++){
        // wait for any rank finish to send another job
        MPI_Recv(&jobnumE, 1, MPI_INT,\
                MPI_ANY_SOURCE, 0, comm, &stat);
        MPI_Recv(&configs[jobnumE], 1, MPI_Config,\
                stat.MPI_SOURCE, 1, comm,\
                MPI_STATUS_IGNORE);

        configs[jobnum].jobnumE = jobnum;
        MPI_Send(&configs[jobnum], 1, MPI_Config,\
                stat.MPI_SOURCE, 0, comm);
    }

    // for the last ones just recv energies and\
    send a -1 to says that its over
    for(dest=1; dest<Nranks; dest++){
        MPI_Recv(&jobnumE, 1, MPI_INT,\
                MPI_ANY_SOURCE, 0, comm, &stat);
        MPI_Recv(&configs[jobnumE], 1, MPI_Config,\
                stat.MPI_SOURCE, 1, comm,\
                MPI_STATUS_IGNORE);

        configs[jobnum].jobnumE = -1;
        MPI_Send(&configs[jobnum], 1, MPI_Config,\
                stat.MPI_SOURCE, 0, comm);
    }
}

// worker process
else{
    while(1){
        MPI_Recv(&configs[0], 1, MPI_Config,\
                0, 0, comm, MPI_STATUS_IGNORE);
        if(configs[0].jobnumE>=0) {
            // update Nranks ensembles at the same time,\
            one per rank
            for(k=0; k<Kmax; k++){
                update(configs[0].x, configs[0].y,\
                        allT[configs[0].jobnumE]);
            }

            // measure the energy and send back to the root
            configs[0].E=ener(configs[0].x, configs[0].y);
            MPI_Send(&configs[0].jobnumE, 1, MPI_INT,\
                    0, 0, comm);
            MPI_Send(&configs[0], 1, MPI_Config,\
                    0, 1, comm);
        }
        else{
            break;
        }
    }
}

```

Results

